CSE 158/258, MGTA 461, DSC 256, Fall 2024: Homework 1

Instructions

Please submit your solution **by Monday Oct 14.** Submissions should be made on **gradescope**. Please complete homework **individually**.

You should submit two files:

`answers_hw1.txt` should contain a python dictionary containing your answers to each question. Its format should be like the following:

> { "Q1": 1.5, "Q2": [3,5,17,8], "Q2": "b", (etc.) }

The provided code stub demonstrates how to prepare your answers and includes an answer template for each question.

`homework1.py` A python file containing working code for your solutions. The autograder *will not execute your code*; this file is required so that we can assign partial grades in the event of incorrect solutions, check for plagiarism, etc. Your solution should **clearly document which sections correspond to each question and answer**. We may occasionally run code to confirm that your outputs match submitted answers, so **please ensure that your code generates the submitted answers.**

You will need the following files:

Homework 1 stub : `https://cseweb.ucsd.edu/classes/fa24/cse258-b/stubs/`

GoodReads Young Adult Reviews : `https://cseweb.ucsd.edu/classes/fa24/cse258-b/data/young_adult_10000.json.gz`

Beer Reviews : `https://cseweb.ucsd.edu/classes/fa24/cse258-b/data/beer_50000.json`

The above are *json* formatted datasets. Code to read them is included in the stub.

Further code examples for regression and classification are available on the class and textbook webpages. Executing the code requires a working install of Python 2.7 or Python 3 with the scipy packages installed.

Each question is worth one mark unless otherwise specified.

Tasks — Regression:

First, using the (GoodReads) *book review* data, let's see whether ratings can be predicted as a function of a few simple features from a review.

1. Train a simple predictor that estimates a rating from the number of times the exclamation mark (!) symbol is used in the review (i.e., `review.count('!')`):

$$\text{star rating} \simeq \theta_0 + \theta_1 \times [\text{number of '!' characters}]$$

   Report the values θ_0 and θ_1, and the Mean Squared Error of your predictor (on the entire dataset).

2. Re-train your predictor so as to include a second feature based on the length (in characters), i.e.,

$$\text{star rating} \simeq \theta_0 + \theta_1 \times [\text{length}] + \theta_2 \times [\text{number of '!' characters}]$$

   Report the values of θ_0, θ_1, and θ_2, along with the MSE of the new model.

3. Train a model that fits a polynomial function to estimate ratings based on our '!' feature. I.e.,

$$\text{star rating} \simeq \theta_0 + \theta_1 \times [\text{number of !}] + \theta_2 \times [\text{number of !}]^2 + \theta_3 \times [\text{number of !}]^3$$

   Fit polynomials from degree one to five (i.e., including up to $[\text{number of !}]^5$) and report the MSE of each.

4. Repeat the above question, but this time split the data into a training and test set. You should split the data into 50%/50% train/test fractions. The *first half* of the data should be in the training set and the *second half* in the test set.[1] Report the MSE of each model on the *test* set.

5. Given a trivial predictor, i.e., $y = \theta_0$, what is the best possible predictor (i.e., value of θ_0) in terms of the *Mean Absolute Error* (MAE), i.e., $\frac{1}{N} \sum_i |y_i - \theta_0|$? For your answer report the MAE of your predictor on the test set from the previous question.

Tasks — Classification:

In this question, using the *beer review* data, we'll try to predict user gender based on users' beer reviews. Load the 50,000 beer review dataset, discarding any entries that don't include a specified gender.

6. Fit a logistic regressor that estimates gender from the number of '!' characters, i.e.,

$$p(\text{gender is female}) \simeq \sigma(\theta_0 + \theta_1 \times [\text{number of !}])$$

Report the number of True Positives, True Negatives, False Positives, False Negatives, and the Balanced Error Rate of the predictor (your answer should be a list of length 5). You may use a logistic regression library with default parameters, e.g. *linear_model.LogisticRegression()* from *sklearn*.

7. Retrain the regressor using the `class_weight='balanced'` option, and report the same error metrics as above.

8. Report the precision@K of your balanced classifier for $K \in [1, 10, 100, 1000, 10000]$ (your answer should be a list of five precision values).

---

[1]Although this isn't best practice (compared to random splitting), it is easier for autograding!